

Domain narrative & design goals

The chess-club database is **read-heavy**: new games arrive slowly, while players query histories, leaderboards and analytics many times a day.

Design therefore follows two main principles:

1. **Query-pattern-first**: each required question must be answerable by *one Redis operation* or, at worst, a short script that touches a single key family.
2. **Write-once, read-many**: we pre-aggregate during **Write Events** and store redundant counts rather than scan large lists later.

Note: the assignment restrict our usage of Redis types to only Lists, Sets, Keys and BloomFilters.

Write Events

ID	Event
E1	Initial load (via <code>load_transform.py</code>)
E2	When a new player is added
E3	When a new game is scheduled
E4	When a game record is inserted
E5	When 72 hours passes on a scheduled game

Key Generation Method

Namespace convention

```
<root>:<entity-id>[:subid][:attribute]
```

- `player` and `game` are the only entity roots.
- A trailing plural (`games` , `opponents`) always holds a **collection** (List / Set).
- `global:` and `analytics:` roots hold cross-player aggregates.

Keys

Player

#	Key pattern	Redis type	Value contents	Written / updated	Read by
P1	player:{pid}:email	string	e-mail address	E1, E2	
P2	player:{pid}:wins	string (int)	win counter	E1, E2, E4	filtered_fof
P3	player:{pid}:losses	string (int)	loss counter	E1, E2, E4	
P4	player:{pid}:games	list of strings	game_id s	E1, E2, E4	player_match_history(pid)
P5	player:{pid}:games-set	set of strings	game_id s	E1, E2, E4	games_against_opponent(pid, oid)
P6	player:{pid}:scheduled	list of strings	future game_id s	E1, E2, E3, E5	future_games(pid)
P7	player:{pid}:opponents	Set of strings	opponent PIDs	E1, E2, E4	friends_of_friends(pid)
P8	player:{pid}:seq	bloom filter	3-move string by PID	E1, E2, E4	player_seq()
P9	player:{pid}:openings:{eco}	string (int)	player ECO counter	E1, E2, E4	player_most_freq_opening(pid)
P10	player:{pid}:most_freq_opening	string	most used opening	E1, E2, E4	player_most_freq_opening(pid)

Game

#	Key pattern	Redis type	Value contents	Written / updated	Read by
G1	game:{gid}:winner	string	winner color	E1, E4	
G2	game:{gid}:victory_status	string	how winner won	E1, E4	
G3	game:{gid}:number_of_turns	string (int)	turns counter	E1, E4	
G4	game:{gid}:white_player_id	string	white player PID	E1, E4	
G5	game:{gid}:black_player_id	string	black player PID	E1, E4	
G6	game:{gid}:opening_eco	string	opening move	E1, E4	
G7	game:{gid}:moves	list of strings	game moves	E1, E4	three-move extraction (loader)

Analytics

#	Key pattern	Redis type	Value contents	Written / updated	Read by
A1	analytics:shortest_game	string	current shortest game	E1, E4	shortest_game()
A2	analytics:check:{gid}	string (int)	“+” count	E1, E4	check_counts()
A3	analytics:top_wins	list of strings	sorted list of joined pid:count s	E1, E4	top_10()
A4	analytics:bottom_losses	list of strings	sorted list of joined pid:count s	E1, E4	bottom_10()

#	Key pattern	Redis type	Value contents	Written / updated	Read by
A5	analytics:most_freq_opening	string	current leader ECO	E1, E4	most_freq_opening()
A6	analytics:most_freq_opening_count	string (int)	its count (avoids extra GET)	E1, E4	most_freq_opening()
A7	analytics:most_common_seq	string	most used 3-move string	E1, E4	most_common_seq()
A8	analytics:least_common_seq	string	least used 3-move string	E1, E4	least_common_seq()
A9	analytics:least_common_seq_count	string (int)	current minimum count	E1, E4	least_common_seq()

Global

#	Key pattern	Redis type	Value contents	Written / updated	Read by
GA1	global:seq	bloom filter	every 3-move string seen	E1, E4	global_seq(seq)
GA2	global:seq:{seq}:count	string (int)	every 3-move string seen	E1, E4	analytics queries
GA3	global:opening:{eco}:count	string (int)	total games that used this ECO	E1, E4	most_freq_opening()
GA4	global:players:emails	set of strings	player_email	E1, E2	in_league(email)

Write-up

Player Queries

player_match_history(player_id)

- **Key:** `player:{pid}:games`
- **Value:** List of every game-id the player has participated in (most-recent first).
- **Example:** `["g17", "g31", "g54"]`
- **Update:**
 - **E1:** initial load pushes all historical game-ids (oldest → newest) so the final list is in recency order.
 - **E2:** create an empty list for a new player.
 - **E4:** for each finished game, `LPUSH` the new `gid` onto each participant's list.
- **Read:** Use `player:{pid}:games` to return the full history in one operation.
- **Cost:**
 - Two `LPUSH` calls per game record (one for each player), and one `LRANGE` per query.
 - Each operation is $O(1)$ for push and $O(N)$ to scan the list of length N , but it's a single Redis call.

future_games(player_id)

- **Key:** `player:{pid}:scheduled`
- **Value:** List of upcoming `game_id` s in the order they were scheduled.
- **Example:** `["g88", "g90"]`
- **Update:**
 - **E1:** initial load builds any pre-scheduled games.
 - **E2:** create an empty list.
 - **E3:** `RPUSH` the new `gid` when a match is scheduled.
 - **E5:** external timer/job invokes removal after 72 h removing older games for each player.
- **Read:** Use `player:{pid}:scheduled` to return the list in one go.
- **Cost:**
 - One `RPUSH` per player when scheduling.
 - One `LREM` per player when expiring.
 - One `LRANGE` per query.

in_league(player_email)

- **Key:** `global:players:emails`
- **Value:** Set of every registered e-mail address.
- **Example:** `{"bob@club.org", "alice@club.org"}`
- **Update:**
 - **E1:** during initial load, `SADD` each player's email.
 - **E2:** when adding a new player, `SADD` their email.
- **Read:** Use `SISMEMBER` with `global:players:emails` to return a boolean in one operation.
- **Cost:**
 - One `SADD` per player write.
 - one `SISMEMBER` per lookup.
 - Both $O(1)$.

games_against_opponent(pid, oid)

- **Keys:**
 - `player:{pid}:games-set`
 - `player:{oid}:games-set`
- **Value:** Each is the full set of games per player.

- **Example:** Intersecting {"g31", "g54"} with {"g31", "g88"} → {"g31"}
- **Update:**
 - **E1:** initial load does SADD for each historical game.
 - **E2:** create an empty set.
 - **E4:** on game insert, two SADD calls (one per participant).
- **Read:** Use SINTER with the two keys to get the answer in one operation.
- **Cost:**
 - Two SADD per game record
 - one SINTER per query.
 - O(N).

player_most_freq_opening(pid)

- **Keys:**
 - player:{uid}:openings:{eco}
 - player:{uid}:most_freq_opening
 - **Value:** A string representing the most used opening by this player.
 - **Example:**
 - player:{pid}:openings:{eco} → 13
 - player:{pid}:most_freq_opening → f3
- **Update:**
 - **E2:** Init to zero.
 - **E1/E4:**

```
SET OR INCR player:{pid}:openings:{eco} → n
GET player:{pid}:most_freq_opening → fav
GET player:{pid}:openings:{fav} (skip if None) → fcnt
IF fav is None OR n > fcnt:
    SET player:{pid}:most_freq_opening {eco}
```

- **Read:** Use GET on player:{pid}:most_freq_opening to get it in one operation.
- **Cost**
 - Reading is one operation.
 - Writing is ~4 operations.

Analytics Queries

most_freq_opening()

- **Key:** analytics:most_freq_opening
- **Value:** A string representing the most used ECO code ever.
- **Example:** analytics:most_freq_opening → C65
- **Update:**
 - **E2:** initialize to empty or "" .
 - **E1/E4:**

```
SET / INCR global:opening:{eco}:count → n
GET analytics:most_freq_opening_count → best # None on first game
IF best is None OR n > best THEN
    SET analytics:most_freq_opening {eco}
    SET analytics:most_freq_opening_count n
```

- **Read:** A simple GET on the key to read it in one query.
- **Cost:**
 - On write, ?.
 - On read, one GET per query O(1).

most_common_seq()

- **Key:** analytics:most_common_seq
- **Value:** The 3-move sequence with the **highest** global count.
- **Example:** analytics:most_common_seq → "e4 e5 Nf3"
- **Update:**
 - **E1:**
 - after writing every global:seq:{seq}:count , run a one-off Python loop that scans those keys
 - finds the max count
 - SET the analytics:most_common_seq key.
 - **E2:** initialize to the empty string "" .
 - **E4:** for every 3-move span extracted from a new game:
 1. INCR global:seq:{seq}:count → newCount
 2. GET analytics:most_common_seq → bestSeq
 3. GET global:seq:{bestSeq}:count → bestCount
 4. If newCount > bestCount , SET analytics:most_common_seq {seq} .

- **Read:** One `GET` operation on `analytics:most_common_seq`.
- **Cost**
 - **Writes:**
 - One `INCR` per span $O(1)$
 - Plus `GET` and `SET` when a new maximum appears.
 - **Reads:** one `GET` per query $O(1)$

`least_common_seq()` (Unstable and I don't like it)

- **Key:** `analytics:least_common_seq`
- **Value:** The 3-move sequence with the lowest positive global count.
- **Example:** `analytics:least_common_seq` \rightarrow "a2 a3 h6"
- **Update:**
 - **E1:**
 - after writing every `global:seq:{seq}:count`, run a one-off Python loop that scans those keys
 - finds the minimum non-zero count
 - `SET` the `analytics:least_common_seq` key.
 - **E2:** initialize to the empty string "".
 - **E4:** for every 3-move span extracted from a new game:

```
new_count = INCR global:seq:{seq}:count

least_seq = GET analytics:least_common_seq
least_count = GET analytics:least_common_seq_count

# first ever value or a strictly smaller count
if least_count is None or new_count < int(least_count):
    MSET analytics:least_common_seq {seq} \
        analytics:least_common_seq_count {new_count}

# brand-new sequence (new_count == 1) is always the new minimum
elif new_count == 1:
    MSET analytics:least_common_seq {seq} \ analytics:least_common_seq_count 1
```

- **Read:** One `GET` operation on `analytics:least_common_seq`.
- **Cost**
 - **Writes:**
 - One `INCR` per span $O(1)$

- Plus GET and SET when a new minimum appears.
- **Reads:** one GET per query $O(1)$

Leaderboard Queries

top_10()

- **Key family:** analytics:top_wins
- **Value:** List of at most 10 player_id s, sorted descending by win count.
- **Example:** ["42", "17", ..., "5"]
 - **E1:** load-time script sorts all players by wins, then RPUSH the first 10 into analytics:top_wins.
 - **E2:** RPUSH new PID at the tail with a 0 count.
 - **E4:** on each win:

```

wins = INCR player:{pid}:wins           # W1
oldTuple = f"{pid}:{wins-1}"
LREM analytics:top_wins 0 oldTuple      # W2  remove stale copy

```

ids = LRANGE analytics:top_wins 0 -1 # R1 single read

find insert position (at most 10 comparisons)

```

pos = first i where wins > int(ids[i].split(':')[1])
newTuple = f"{pid}:{wins}"

```

if pos is None: # wins is now the smallest

```
RPUSH analytics:top_wins newTuple # W3
```

else:

```
LINSERT analytics:top_wins BEFORE ids[pos] newTuple # W3
```

if LLEN analytics:top_wins > 10:

```
RPOP analytics:top_wins # W4 trim list
```

- **Read:** Use LRANGE on analytics:top_wins and get it in one query.

- **Cost:**
 - Reading is 1 op (LRange)
 - Writing ≤ 4 ops (INCR , LREM , insert, optional RPOP).

bottom_10()

- **Key:** analytics:bottom_losses
- **Value:** List of the **10 highest-loss players**,.
- **Example:** ["99", "8", ..., "23"]
- **Update**
 - **E1:** build initial list.
 - **E2:** RPush new PID with 0.
 - **E4:**

```

losses = INCR player:{pid}:losses                # W1

oldTuple = f"{pid}:{losses-1}"
LREM analytics:bottom_losses 0 oldTuple          # W2

ids = LRange analytics:bottom_losses 0 -1        # R1

pos = first i where losses > int(ids[i].split(':')[1])
newTuple = f"{pid}:{losses}"

if pos is None:
    RPush analytics:bottom_losses newTuple        # W3
else:
    LINSERT analytics:bottom_losses BEFORE ids[pos] newTuple # W3

if LLEN analytics:bottom_losses > 10:
    RPOP analytics:bottom_losses                  # W4

```

- **Read:** Use LRange on analytics:bottom_losses and get it in one query.
- **Cost:**
 - Reading is 1 op (LRange)
 - Writing ≤ 4 ops (INCR , LREM , insert, optional RPOP).

Game Queries

`player_seq(user, seq)` (Incorrect)

- **Key:** `player:{pid}:seq`
- **Value:** ...
- **Example:** Filter `"e4 e5 Nf3"` .
- **Update:** on E4, for each game's move list, take every consecutive span of three half-moves and `BF.ADD` to the filter.
- **Read:** Use `BF.EXISTS` with `player:{pid}:seq` in one call.
- **Cost:**
 - One `BF.EXISTS` per read query $O(1)$.
 - Writes proportional to the number of moves per game.

`global_seq(seq)` (Incorrect)

- **Key:** `global:seq`
- **Value:** ...
- **Example:** `BF.EXISTS global:seq "e4 e5 Nf3"` \rightarrow 1 (present) or 0 (absent)
- **Update**
 - **E1:** during initial load, for each three-move span do `BF.ADD global:seq seq` .
 - **E4:** on each new game, for each three-move span do `BF.ADD global:seq seq` .
- **Read:** One `BF.EXISTS` on `global:seq "{seq}"`
- **Cost**
 - One `BF.EXISTS` per query $O(1)$.
 - Writes: one `BF.ADD` per three-move span.

Analytics Queries

`shortest_game()`

- **Key:** `analytics:shortest_game`
- **Value:** Game-id of the record holder for fewest turns.
- **Example:** `"g17"`
- **Update:**
 - On **E1**, set to the shortest from the load.
 - On **E4**, compare the new game's turn count and `SET` if lower.
- **Read:** Simple `GET` on `analytics:shortest_game` .
- **Cost:**

- One comparison + possible write per game.
- One $O(1)$ read per query.

check_counts(game_id)

- **Key:** analytics:check:{gid}
- **Value:** Integer count of “+” symbols in that game’s moves.
- **Example:** 5
- **Update:** On **E1/E4**, count the occurrences and SET the analytics:check:{gid} key.
- **Read:** Use GET with analytics:check:{gid}.
- **Cost:**
 - One write per game $O(1)$
 - One $O(1)$ read per query.

Graph Queries

friends_of_friends(user)

- **Key:** player:{pid}:opponents
- **Value:** Set of every opponent the user has faced.
- **Example:** {"17", "99", "103"}
- **Update:**
 - On **E1**, populate from history
 - On **E2**, create empty set
 - On **E4**, two SADD calls (one per player).
- **Read:** To be written by the person working on Graph queries.
- **Cost:** To be written by the person working on Graph queries.

filtered_friends_of_friends(user)

- **Keys:** Same FoF set plus player:{pid}:wins counters.
- **Value:** FoF PIDs who have more wins than the user.
- **Example:** {"17", "103"} after filtering.
- **Update:** No extra writes beyond wins increments on **E4**.
- **Read:** To be written by the person working on Graph queries.
- **Cost:** To be written by the person working on Graph queries.

`largest_connected_component()`

- **Key fabric:** All `player:{pid}:opponents` sets form the graph.
- **Value:** Computed on demand.
- **Update:** Only the usual `SADD` on **E4**.
- **Read:** To be written by the person working on Graph queries.
- **Cost:** To be written by the person working on Graph queries.